


Arabic Handwritten Characters Recognition using Convolutional Neural Network

CCAI 435 | Deep Learning Course

Dataset Description

The Arabic Handwritten Characters dataset, introduced in a paper by El-Sawy et al., contains all Arabic characters handwritten by 60 participants of varying ages. The dataset is divided into two splits: a training set with 13,440 images and a testing set with 3,360 images all have the same size (32,32). Each character's label is linked to the corresponding image name, which represents the character's letter name. The following figures show the size of dataset and sample images from the dataset:

```
Train Data shape (13440, 32, 32, 3)
Train Labels shape (13440,)
Test Data shape (3360, 32, 32, 3)
Test Labels shape (3360,)
```





Neural Network Architecture & Results

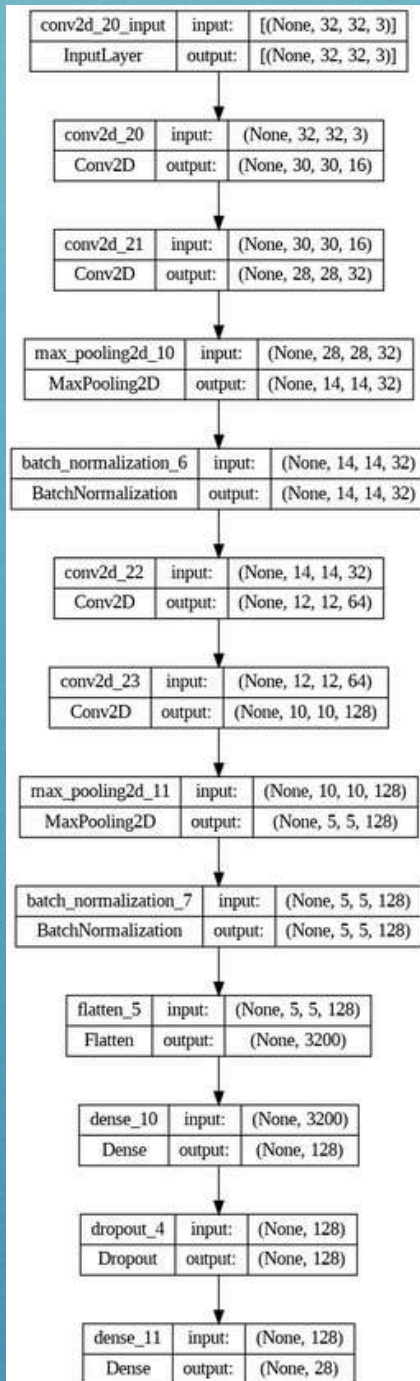
Best Neural Network Architecture

We have constructed a neural network architecture with a multi-layer convolutional neural network (CNN) implemented using TensorFlow's Keras. Below, we will mention the main features of our network:

- **12 layers in total, including:**
- **Feature Extraction:**
 - Conv2D layers with increasing filters (16, 32, 64, 128), all using a filter size of (3,3) and stride of 1
 - MaxPooling and Batch Normalization for regularization and dimension reduction
- **Classification:**
 - Flatten layer to transition from 2D to 1D
 - Dense layers for classification:
 - First layer: 128 units with L1 and L2 regularization ($\lambda_1=0.0001, \lambda_2=0.0001$)
 - Second layer: 28 units with softmax activation
 - Dropout at 50% to prevent overfitting

Layer (type)	Output Shape	Param #
conv2d_20 (Conv2D)	(None, 30, 30, 16)	448
conv2d_21 (Conv2D)	(None, 28, 28, 32)	4640
max_pooling2d_10 (MaxPooling2D)	(None, 14, 14, 32)	0
batch_normalization_6 (Batch Normalization)	(None, 14, 14, 32)	128
conv2d_22 (Conv2D)	(None, 12, 12, 64)	18496
conv2d_23 (Conv2D)	(None, 10, 10, 128)	73856
max_pooling2d_11 (MaxPooling2D)	(None, 5, 5, 128)	0
batch_normalization_7 (Batch Normalization)	(None, 5, 5, 128)	512
flatten_5 (Flatten)	(None, 3200)	0
dense_10 (Dense)	(None, 128)	409728
dropout_4 (Dropout)	(None, 128)	0
dense_11 (Dense)	(None, 28)	3612

Best Neural Network Architecture



```
1 his_6 = model_6.fit(train, train_labels, epochs=10, validation_split=0.2, batch_size=32)
```

Epoch 1/10

336/336 [=====] - 5s 7ms/step - loss: 2.9316 - accuracy: 0.3899 - val_loss: 7.3247 - val_accuracy: 0.0368

Epoch 2/10

336/336 [=====] - 2s 6ms/step - loss: 1.5677 - accuracy: 0.7397 - val_loss: 2.0505 - val_accuracy: 0.5945

Epoch 3/10

336/336 [=====] - 2s 7ms/step - loss: 1.1559 - accuracy: 0.8453 - val_loss: 1.5237 - val_accuracy: 0.7165

Epoch 4/10

336/336 [=====] - 2s 7ms/step - loss: 0.9503 - accuracy: 0.8922 - val_loss: 1.0590 - val_accuracy: 0.8423

Epoch 5/10

336/336 [=====] - 2s 6ms/step - loss: 0.8267 - accuracy: 0.9115 - val_loss: 1.1662 - val_accuracy: 0.8006

Epoch 6/10

336/336 [=====] - 2s 6ms/step - loss: 0.7222 - accuracy: 0.9327 - val_loss: 0.8289 - val_accuracy: 0.8943

Epoch 7/10

336/336 [=====] - 2s 6ms/step - loss: 0.6880 - accuracy: 0.9348 - val_loss: 1.0195 - val_accuracy: 0.8348

Epoch 8/10

336/336 [=====] - 2s 6ms/step - loss: 0.6662 - accuracy: 0.9381 - val_loss: 0.7760 - val_accuracy: 0.9077

Epoch 9/10

336/336 [=====] - 2s 6ms/step - loss: 0.6487 - accuracy: 0.9423 - val_loss: 0.6551 - val_accuracy: 0.9457

Epoch 10/10

336/336 [=====] - 3s 7ms/step - loss: 0.6111 - accuracy: 0.9523 - val_loss: 0.6330 - val_accuracy: 0.9505

```
1 loss, acc = model_6.evaluate(test, test_labels)
```

105/105 [=====] - 0s 3ms/step - loss: 0.6352 - accuracy: 0.9482

Best Neural Network Training and Results

- **Configuration:**

Loss: Sparse categorical crossentropy

Optimizer: Adam

Metric: Accuracy


- **Training:**

Epochs: 10

Validation split: 20%

Batch size: 32

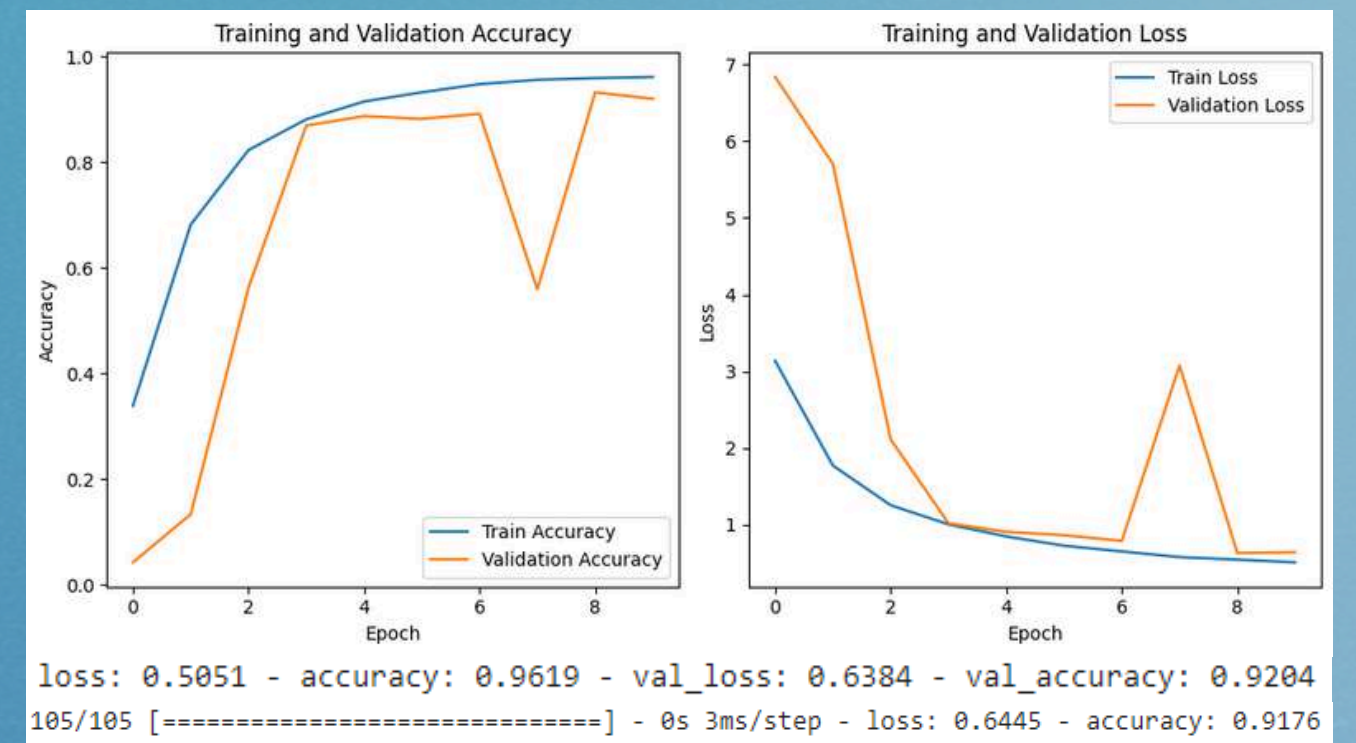

- **Results**



loss: 0.6352 - accuracy: 0.9482

Our Experiment

We trained our final model using over 6 different architectures and hyperparameters. In the process of training, we have had issues like overfitting and low validation accuracy. We have tried to overcome these challenges through dropout and regularization techniques. The following figure shows a sample of our trail to reach the best-performing model architecture and the most recent one for the dataset.





Reference

- Elsayy, Ahmed & Loey, Mohamed & El-Bakry, Hazem. (2017). Arabic Handwritten Characters Recognition using Convolutional Neural Network. WSEAS TRANSACTIONS on COMPUTER RESEARCH. 5. 11-19.
-